

CSE160: Computer Networks

Lecture #04 – Framing, Error Detection and Correction



2020-09-08

**Professor
Alberto E. Cerpa**



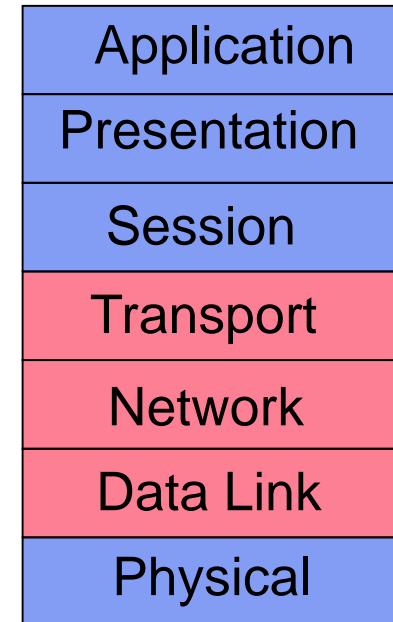
Last Time

- We typically model links in terms of bandwidth and delay, from which we can calculate message latency
- Different media have different properties that affect their performance as links
- We need to encode bits into signals so that we can recover them at the other end of the channel
- Passband modulation allows us to transmit baseline signals in other higher frequencies for fiber and wireless
- Shannon showed us the limit of what is possible to transmit on any channel




This Lecture

- Framing and Error detection and correction
- Focus: How do we send full messages? How do we detect and correct messages that are garbled during transmission?
- The responsibility for doing this cuts across the different layers

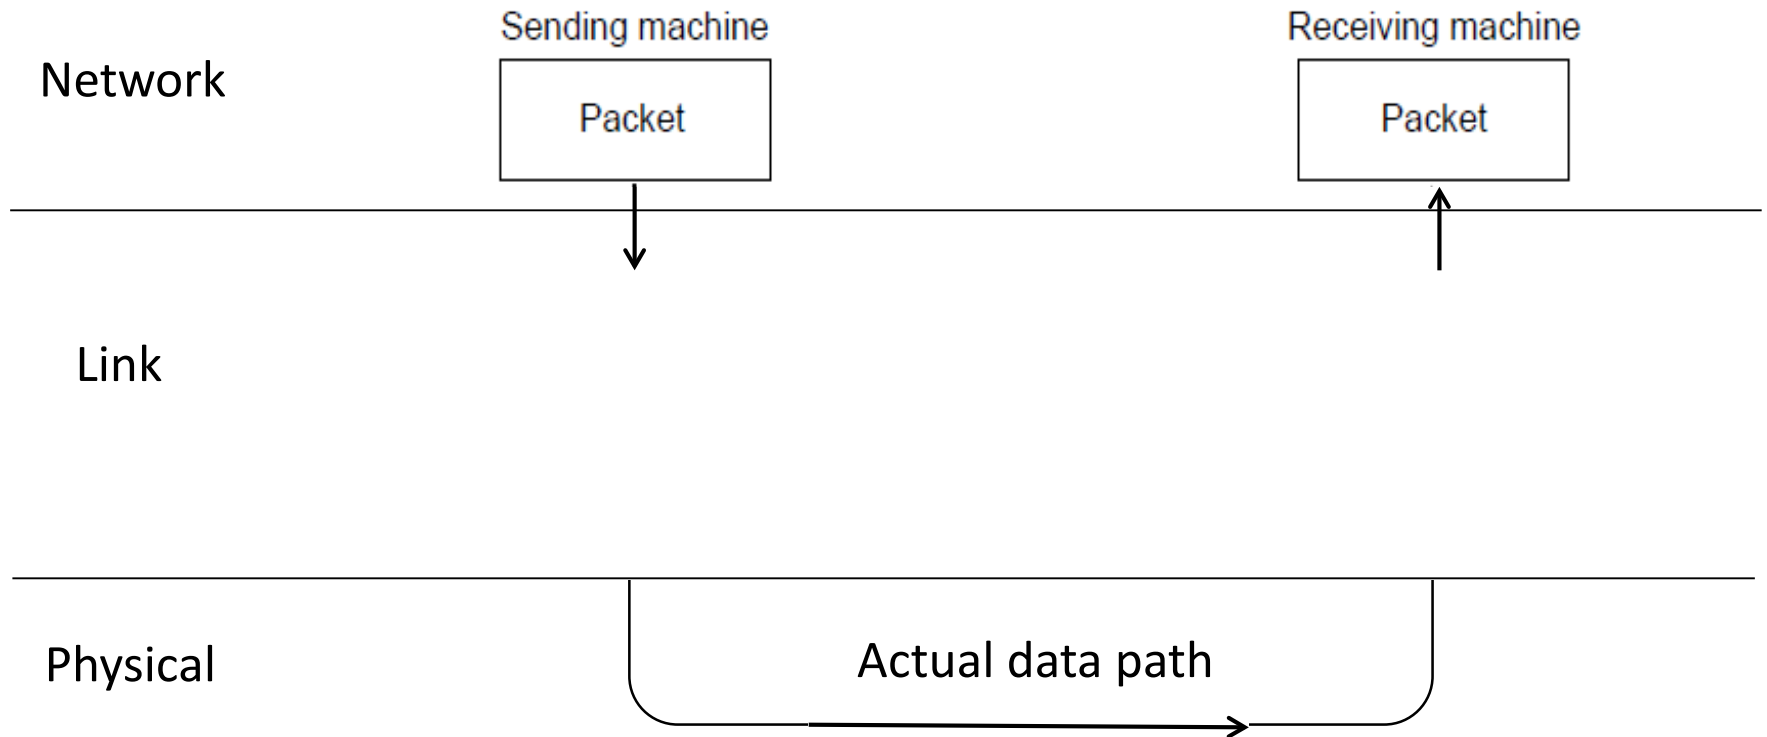


Scope of the Link Layer

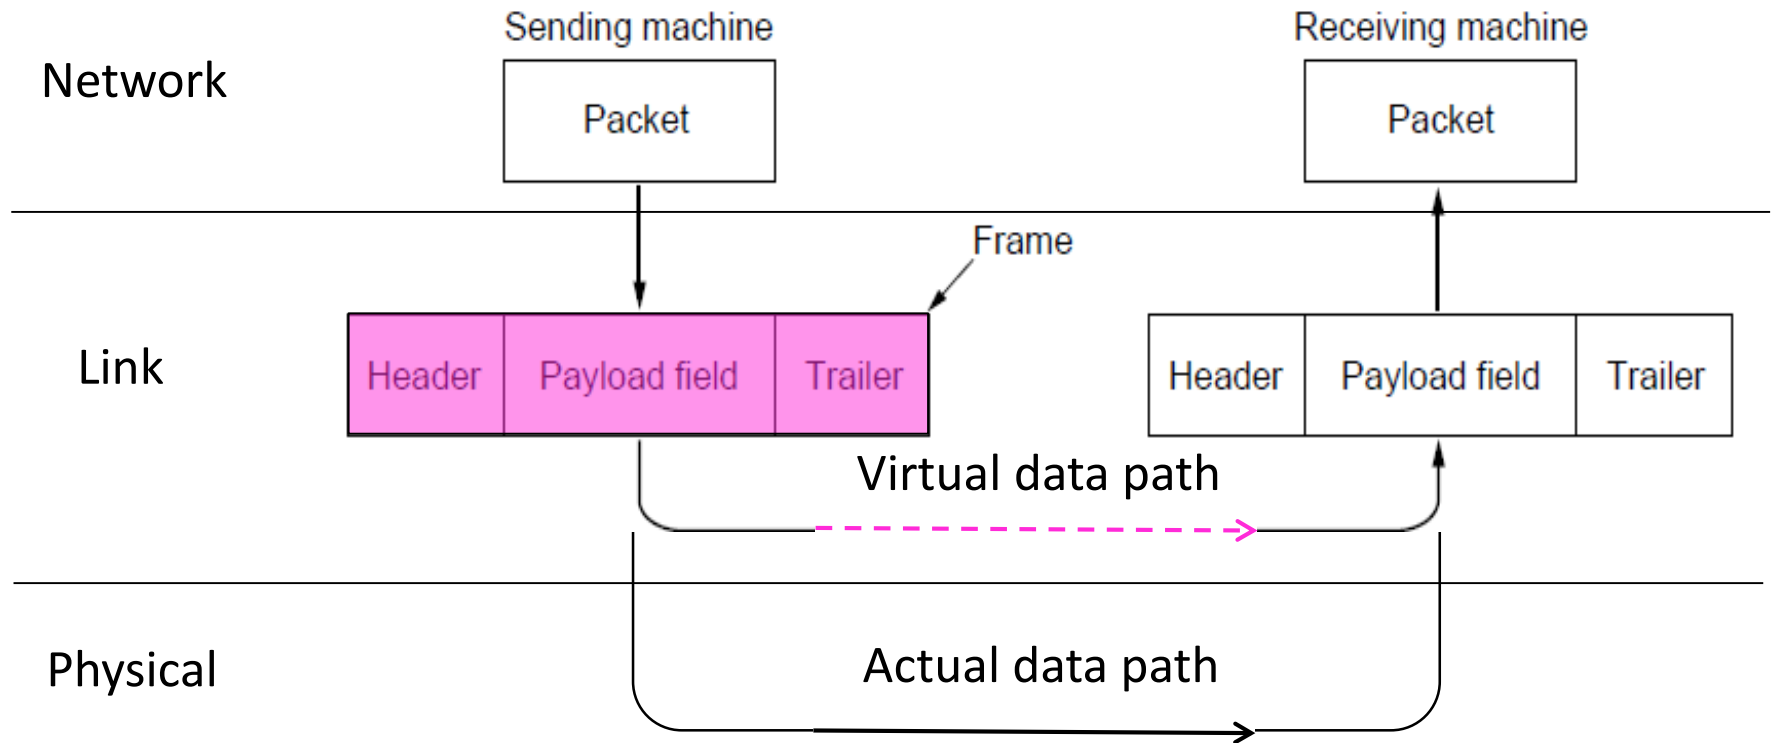
- Concerns how to transfer messages over one or more connected links
 - Messages are frames, of limited size
 - Buildings on the physical layer



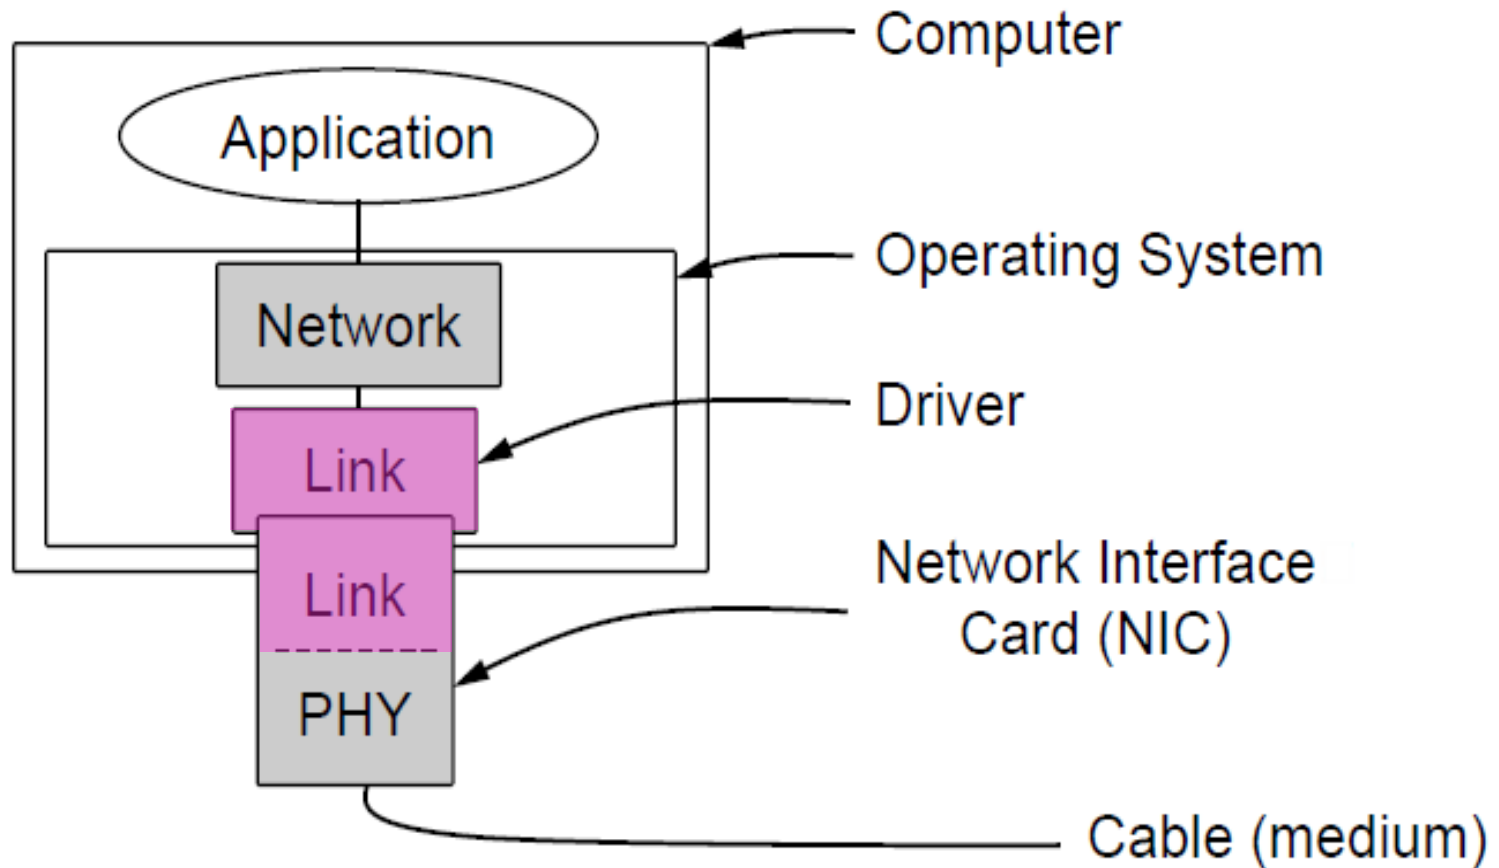
In terms of layers



In terms of layers




Typical Implementation of Layers



Framing

- Need to send message, not just bits
 - The physical layer gives us a stream of bits. How do we interpret it as a sequence of frames?
 - Requires that we synchronize on the start of message reception at the far end of the link



Framing Methods

- We'll look at:
 - Byte count (motivation)
 - Byte stuffing
 - Bit stuffing
- In practice, the physical layer often helps to identify frame boundaries
 - E.g., Ethernet, 802.11



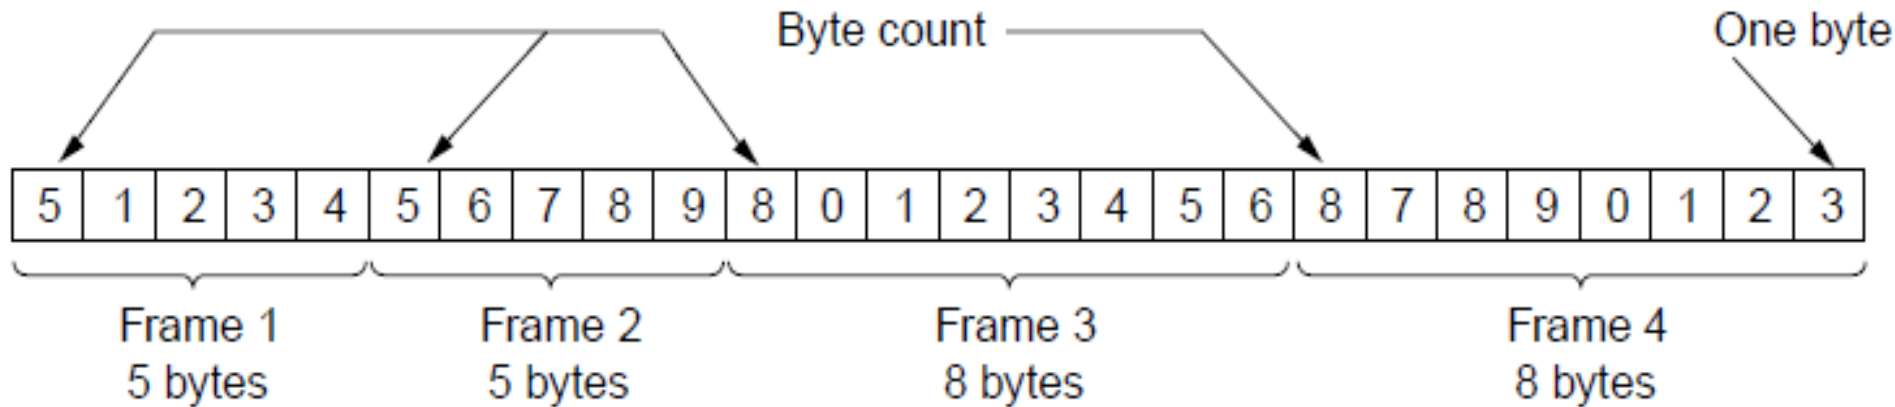
Byte Count

- First try:
 - Let's start each frame with a length field!
 - It's simple, and hopefully good enough ...



Byte Count (2)

- First byte determines the total frame size

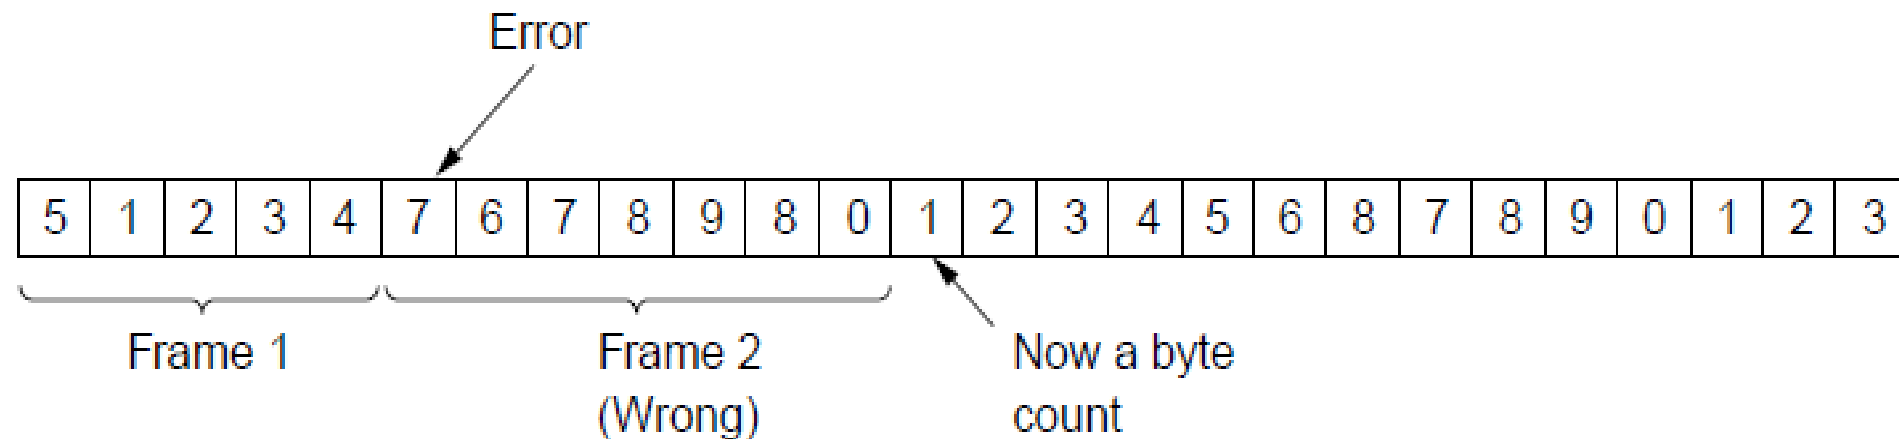


- How well do you think it works?




Byte Count (3)

- Difficult to re-synchronize after framing error
 - Want a way to scan for a start of a frame



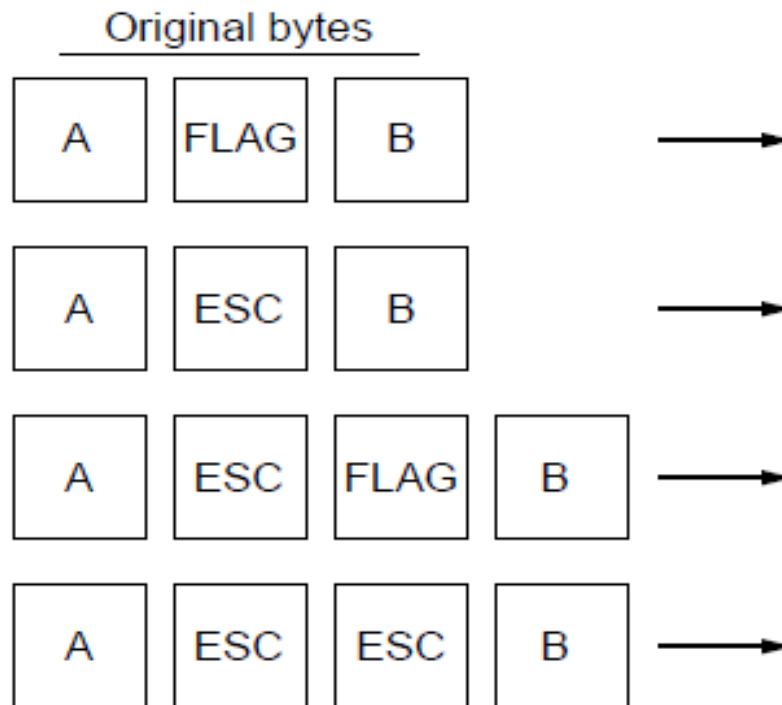
Byte Stuffing

- Better idea:
 - Have a special flag byte value that means start/end of frame
 - Replace (“stuff”) the flag inside the frame with an escape code
 - Complication: have to escape the escape code too!



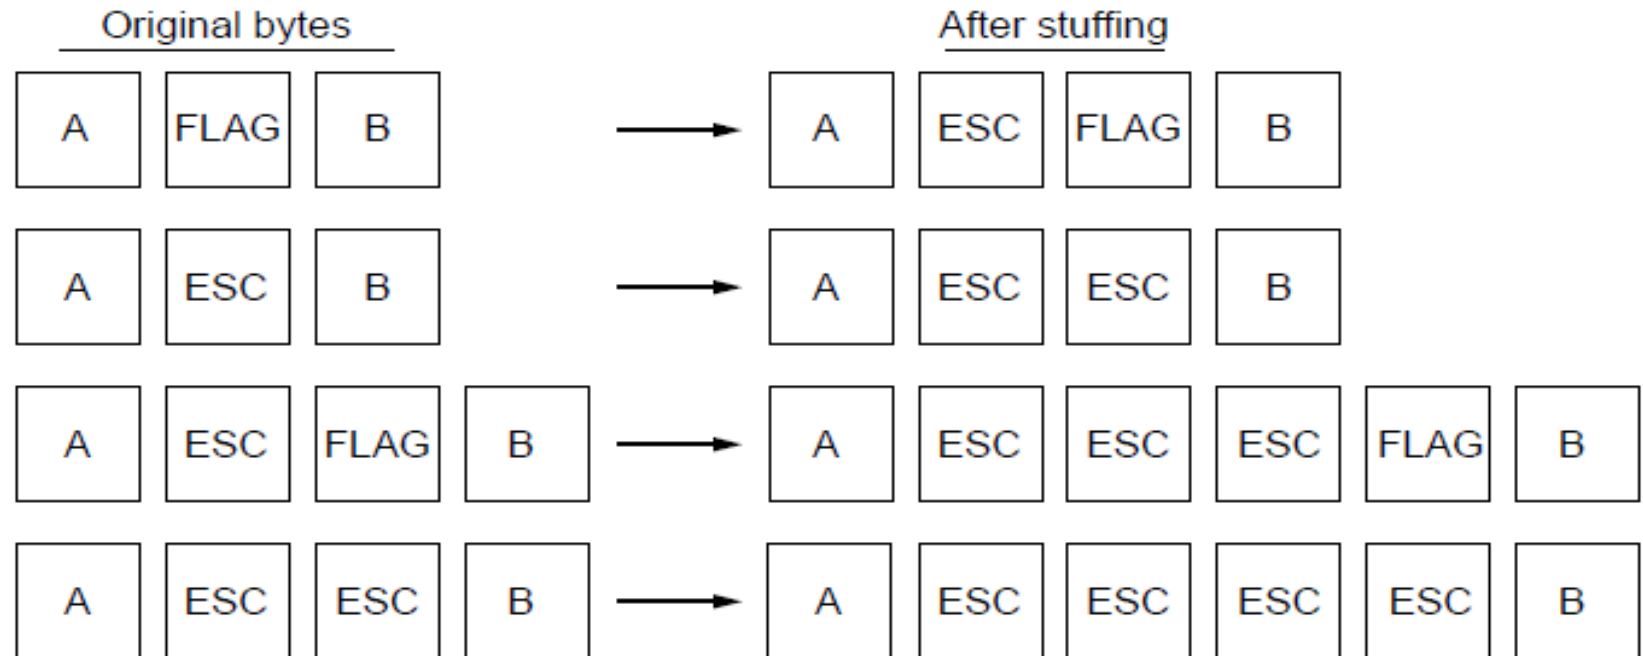
Byte Stuffing (2)

- Rules:
 - Replace each FLAG in data with ESC FLAG
 - Replace each ESC in data with ESC ESC



Byte Stuffing (3)

- Now any unescaped FLAG is the start/end of a frame



Bit Stuffing

- Can stuff at the bit level too
 - Call a flag six consecutive 1s
 - On transmit, after five 1s in the data, insert a 0
 - On receive, a 0 after five 1s is deleted



Bit Stuffing (2)

- Example:

Data bits 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0

Transmitted bits
with stuffing



- So how does it compare with byte stuffing?

Transmitted bits with stuffing

0 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 1 0

Stuffed bits

- It is slightly more efficient but more complicated in practice



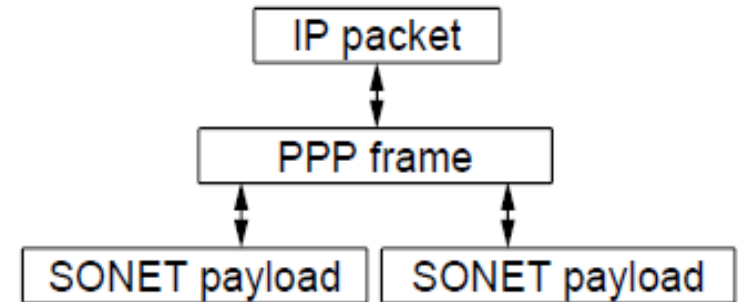
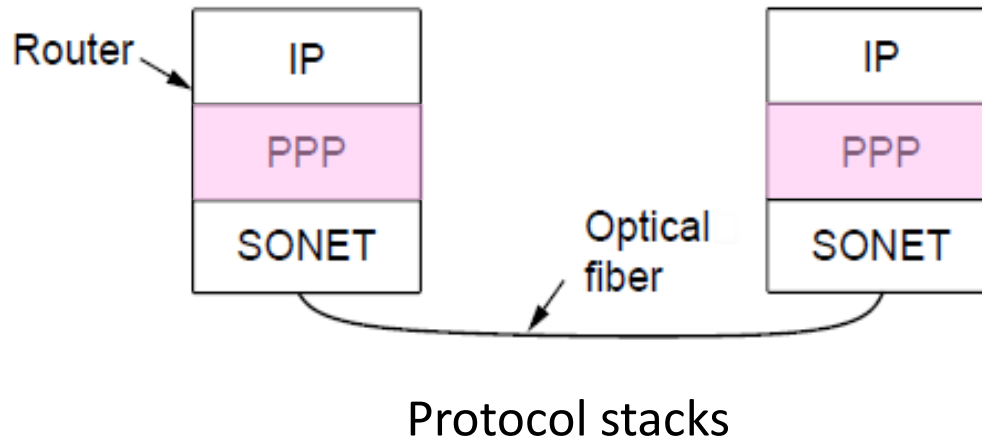
Link Example: PPP over SONET

- PPP is Point-to-Point Protocol
- Widely used for link framing
 - E.g., it is used to frame IP packets that are sent over SONET optical links



Link Example: PPP over SONET (2)

- Think of SONET as a bit stream, and PPP as the framing the carries an IP packet over the link

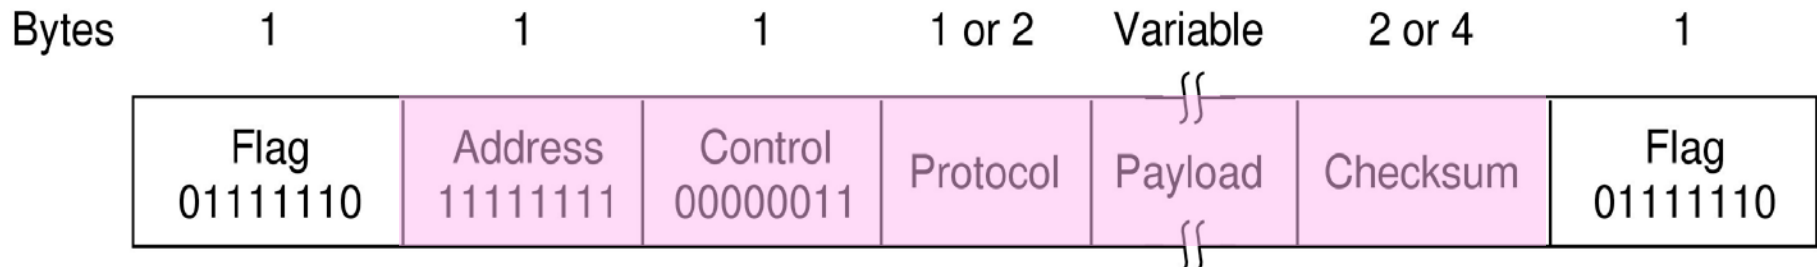


PPP frames may be split over SONET payloads



Link Example: PPP over SONET (3)

- Framing uses Byte stuffing method:
 - FLAG is 0x7E and ESC is 0x7D
 - To stuff (unstuff) a byte, add (remove) ESC (0x7D), and XOR byte with 0x20
 - Removes FLAG from the contents of the frames




Error Coding

- Noise can flip some of the bits we receive
 - We must be able to detect when this occurs!
 - Why?
 - Who needs to detect it? (links, routers, OSs, or apps?)
- What can we do?
 - Detect errors with codes
 - Correct errors with codes
 - Retransmit lost frames ← Later
- Reliability is a concern that cuts across the layers – we'll see it again



Problem – Noise may flip received bits



Approach – Add Redundancy

- Error detection codes allow errors to be recognized
 - Add check bits to the message bits to let some errors be detected
- Error correction codes allow errors to be repaired
 - Add more check bits to let some errors be corrected
- Key issue is now to structure the code to detect many errors with few check bits and modest computation



Motivating Example

- What would be the simplest error detection scheme?
- A simple error detection scheme:
 - Just send two copies. Differences imply errors.
- How good is this code?
 - How many errors can it detect/correct?
 - How many errors will make it fail?



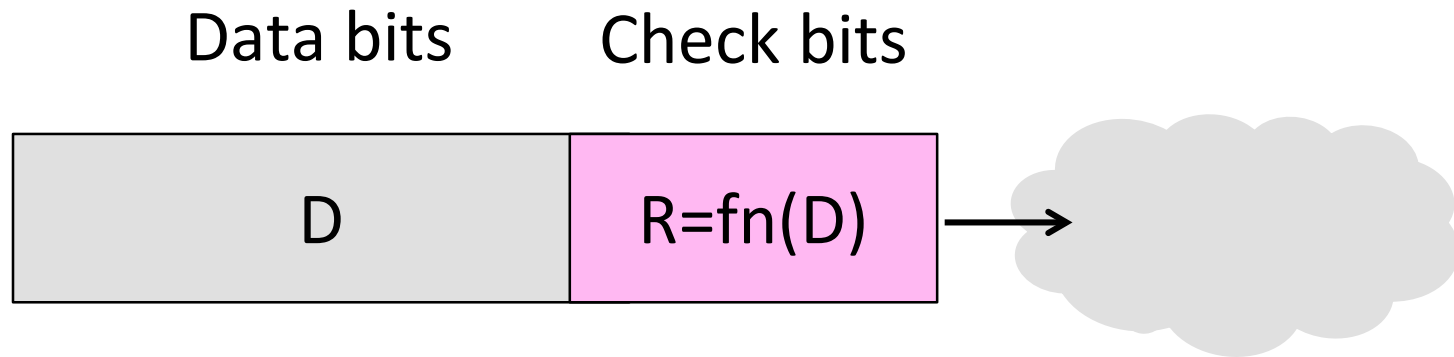
Motivating Example (2)

- We want to handle more errors with less overhead
 - Will look at better codes; they are applied mathematics
 - But, they can't handle all errors
 - And they focus on accidental errors (will look at secure hashes later)
- We will look at basic block codes
 - K bits in, N bits out is a (N,K) code
 - Simple, memoryless mapping



Using Error Codes

- Codeword consists of D data plus R check bits (=systematic block code)

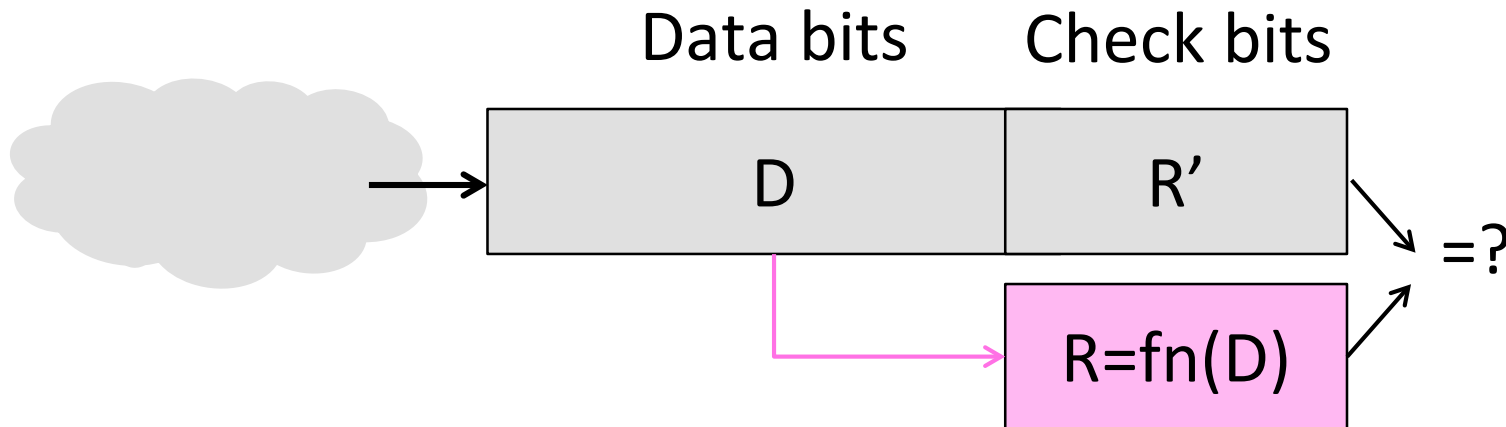


- Sender:
 - Compute R check bits based on the D data bits;
send the codeword of $D+R$ bits




Using Error Codes (2)

- Receiver:
 - Receive $D+R$ bits with unknown errors
 - Recompute R check bits based on the D data bits; error if R doesn't match R'
 - Note the error can be in D , R or both



Intuition for Error Codes

- For D data bits, R check bits:




- Randomly chosen codeword is unlikely to be correct; overhead is low
 - How unlikely?
 - Prob = $1/2^R$



Richard W. Hamming (1915-1998)

- Much early work on codes:
 - “Error Detecting and Error Correcting Codes”, BSTJ, 1950
 - Very easy to read, and very elegant way to explain hamming codes
- See also:
 - “You and Your Research”, 1986



Source: IEEE GHN, © 2019 IEEE



The Hamming Distance

- Errors must not turn one valid codeword into another valid codeword, or we cannot detect/correct them
- Hamming distance of a code is the smallest number of bit differences that turn any one codeword into another (i.e. $(D+R)_1$ to $(D+R)_2$)
 - e.g, code 000 for 0, 111 for 1, Hamming distance is 3
- For code with distance $d = D+1$:
 - D errors can be detected, e.g, 001, 010, 100, 110, 101, 011
- For code with distance $d = 2C+1$:
 - C errors can be corrected, e.g., 001 \rightarrow 000



Error Detection

- Some bits may be received in error due to noise
How do we detect this?
 - Parity
 - Checksums
 - CRCs
- Detection will let us fix the error, for example, by retransmission (later)



Simple Error Detection – Parity Bit

- Have you ever used a modem or a serial connection? Ever heard of parity?
- Start with D bits and add 1 check bit that is the sum of the D bits
 - Sum is modulo 2 or XOR
 - e.g. 0110010 → 01100101
 - Easy to compute as XOR of all input bits
- Will detect an odd number of bit errors
 - But not an even number
- Does not correct any errors
- Hamming distance?



2D Parity

- Add parity row/column to array of bits
- Detects all 1, 2, 3 bit errors, and many errors with >3 bits.
- Corrects all 1 bit errors

		↓	
	0101001	1	
	1101001	0	
	1011110	1	
	0001110	1	
	0110100	1	
	1011111	0	
→	1111011	0	←
		↑	



Checksums

- Idea: sum up the data in N-bit words and send the data along with the sum
 - Widely used in, e.g., TCP/IP/UDP
 - Stronger protection than parity



- Algorithm:
 - checksum is the 1s complement of the 1s complement sum of the data interpreted 16 bits at a time (for 16-bit TCP/UDP checksum)
- 1s complement: flip all bits to make number negative
 - Adding requires carryout to be added back



Internet Checksum Example

- Sending:

1. Arrange data in 16-bit words

2. Put zero in checksum position, add

3. Add any carryover back to get 16 bits

4. Negate (complement) to get sum

$$\begin{array}{r} 0001 \\ \text{f}203 \\ \text{f}4\text{f}5 \\ \text{f}6\text{f}7 \\ + (0000) \\ \hline 2\text{d}\text{d}\text{f}0 \\ \quad \downarrow \\ \text{d}\text{d}\text{f}0 \\ + \quad \quad 2 \\ \hline \text{d}\text{d}\text{f}2 \\ \quad \downarrow \\ 220\text{d} \end{array}$$



Internet Checksum Example (2)

- Receiving:

1. Arrange data in 16-bit words

2. Checksum will be non-zero, add

3. Add any carryover back to get 16 bits

4. Negate (complement) the result and check it is 0

```
0001
f203
f4f5
f6f7
+ 220d
-----
2fffd
  ↓
  fffd
+      2
-----
ffff
  ↓
0000
```



Internet Checksum

- How well does the checksum work?
 - What is the distance of the code?
 - How many errors can detect/correct?
 - Detect: 1, Correct: 0
 - Hamming distance = 2
- What about larger errors?
 - All burst errors up to 16 in length
 - Random errors (random data), with probability $1/2^{16}$



Checksum Code

```
uint16_t
cksum(uint16_t *buf, int count) {
    register uint32_t sum = 0;
    while (count--) {
        sum += *buf++;
        if (sum & 0xFFFF0000) {
            sum &= 0xFFFF;
            sum++;
        }
    }
    return ~(sum & 0xFFFF);
}
```



CRCs (Cyclic Redundancy Check)

- Stronger protection than checksums
 - Used widely in practice, e.g., Ethernet CRC-32
 - Implemented in hardware (XORs and shifts)
- Algorithm: Given n bits of data, generate a k check bits such that the $n+k$ bits that are evenly divisible by a chosen divisor $C(x)$
- Based on mathematics (arithmetic) of finite fields
 - “numbers” correspond to polynomials, use modulo arithmetic
 - e.g, interpret 10011010 as $x^7 + x^4 + x^3 + x^1$



CRC Example

- How do we generate the check sequence?
 - Have our message, e.g., 10011010 ($m=8$)
 - Have the CRC as a divisor polynomial e.g., $C(x) = 1110$ ($x^3 + x^2 + x^1$; $k=3$)
 - Want to make $m + k$ bits divisible by this divisor...
- Send Procedure:
 1. Extend the n data bits with k zeros
 2. Divide by $C(x)$ to find the remainder (ignore quotient)
 3. Adjust k check bits by remainder
- Receive Procedure:
 1. Divide and check for zero remainder



Example – Polynomial Division

Data bits: 10011010


Mesg. bits: 10011010 000

Check bits:

$$C(x) = x^3 + x^2 + 1$$

$$C = 1101$$

$$k = 3$$



Example – Remainder to CRC

- So we see the remainder is 101
- Thus the zero extended message – 101 must be evenly divisible by $C(x)$!
- So perform the subtraction to discover the check bits
 - Subtraction/addition is XOR in module 2 arithmetic
 - E.g., we get $10011010000 - 101 = 10011010101$
 - The check bits are 101
- Finally, message we send is 10011010101



How is $C(x)$ Chosen?

- Mathematical properties:
 - All 1-bit errors if non-zero x^k and x^0 terms
 - All 2-bit errors if $C(x)$ has a factor with at least three terms
 - Any odd number of errors if $C(x)$ has $(x + 1)$ as a factor
 - Any burst error $< k$ bits
 - For random errors, $\text{Prob}(\text{corruption undetected}) = 1/(2^n)$ (e.g. CRC- n with n being 8, 12, 16, 32, etc.)
- There are standardized polynomials of different degree that are known to catch many errors
 - Ethernet CRC-32:
100000100110000010001110110110111
 - Corruption undetected (random) = $2.3 \cdot 10^{-10}$, less than one in a billion



Standard CRC Polynomials

- CRC-8 100000111
 - ITU-T I.432.1 (02/99); ATM HEC, ISDN HEC and others
- CRC-10 11000110011
 - ATM; ITU-T I.610, mobile networks
- CRC-12 110000000111
 - Telecom systems
- CRC-16 1000100000100000
 - X.25, V.41, HDLC FCS, XMODEM, Bluetooth, PACTOR, SD, DigRF, many others
- CRC-32 100000100110000010001110110111
 - ISO 3309 (HDLC), ISO/IEC/IEEE 802-3 (Ethernet), SATA, MPEG2, PKZIP, Gzip, Bzip2, POSIX cksum, PNG, ZMODEM, many others



Error Detection in Practice

- CRC are widely used on links
 - Ethernet, 802.11, ADSL, Cable...
- Checksum used in Internet
 - IP, TCP, UDP, ... but it is weak
- Parity
 - Is little used



Error Correction

- Some bits may be received in error due to noise. How do we fix them?
 - Hamming code
 - Other codes
- And why should we use detection when we can use correction?



Why Error Correction is Hard?

- If we had reliable check bits we could use them to narrow down the position of the error
 - Then correction would be easy
- But error could be in the check bits as well as the data bits!
 - Data might even be correct!




Intuition for Error Correcting Code

- Suppose we construct a code with a Hamming distance of at least 3
 - Need ≥ 3 bit errors to change one valid codeword into another
 - Single bit errors will be closest to a unique valid codeword
- If we assume errors are only 1 bit, we can correct them by mapping an error to the closest valid codeword
 - Works for d errors if $HD \geq 2d + 1$




Intuition (2)

- Visualization of code:



Intuition (3)

- Visualization of code:



Hamming Code

- Gives a method for constructing a code with a distance of 3
 - Uses $n = 2^k - k - 1$, e.g., $n=4$ $k=3$
 - Put check bits in positions p that are powers of 2, starting with position 1
 - Check bit in position p is parity of positions with a p term in their values
- Plus an easy way to correct [soon]



Hamming Code (2)

- Example: data = 0101, 3 check bits
 - 7 bits code, check bit positions 1, 2, 4
 - Check 1 covers positions 1, 3, 5, 7
 - Check 2 covers positions 2, 3, 6, 7
 - Check 4 covers positions 4, 5, 6, 7

0 1 0 0 1 0 1
1 2 3 4 5 6 7

- $P_1 = 0+1+1 = 0$
- $P_2 = 0+0+1 = 1$
- $P_4 = 1+0+1 = 0$



Hamming Code (3)

- To decode:
 - Recompute check bits (with parity sum including the check bit)
 - Arrange as a binary number
 - Value (syndrome) tells error position
 - Value of zero means no error
 - Otherwise, flip bit to correct



Hamming Code (4)

- Example, continued

0 1 0 0 1 0 1
1 2 3 4 5 6 7

- $P_1 = 0+0+1+1 = 0$
- $P_2 = 1+0+0+1 = 0$
- $P_4 = 0+1+0+1 = 0$
- Syndrome = 000 (no error)
- Data = 0 1 0 1



Hamming Code (5)

- Example, continued

0 1 0 0 1 **1** 1
1 2 3 4 5 6 7

- $P_1 = 0+0+1+1 = 0$
- $P_2 = 1+0+\mathbf{1}+1 = \mathbf{1}$
- $P_4 = 0+1+\mathbf{1}+1 = \mathbf{1}$
- Syndrome = **110** \rightarrow 6 (6th position), flip position 6
- Data = 0 1 0 1 (corrected after flip!)



Other Error Correction Codes


- Codes used in practice are much more involved than Hamming
- Convolutional codes
 - Take a stream of data and output a mix of the recent input bits
 - Makes each output bit less fragile
 - Decode using Viterbi algorithm (which can use bit confidence values)



Reed-Solomon Code

- Developed to protect data on magnetic disks
- Used for CDs, Blue-Rays and cable modems too
- Property: $2t$ redundant bits can correct $\leq t$ errors
- Slowly being replaced by low-density parity-check (LDPC) or turbo codes
- Mathematics somewhat more involved ...

Irving S. Reed




Source: IEEE GHN, © 2019 IEEE



Other Codes -- LDPC

- Low Density Parity Check
 - LDPC based on sparse matrices
 - Decoded iteratively using a belief propagation algorithm
 - State of the art today
- Invented by Robert Gallager in 1963 as part of his PhD thesis
 - Promptly forgotten until 1996...



Source: IEEE GHN, © 2019 IEEE



Detection vs. Correction

- Two strategies to correct errors:
 - Detect and retransmit, or Automatic Repeat reQuest. (ARQ)
 - Error correcting codes, or Forward Error Correction (FEC)
- Satellites, real-time media tend to use error correction
- Retransmissions typically at higher levels (Network+)
- Question: Which should we choose?



Retransmissions vs. FEC

- The better option depends on the kind of errors and the cost of recovery
- Example: Message with 1000 bits, with a bit error rate (BER) of 1 in 1000, i.e., $\text{prob}(\text{bit error}) = 0.001$
 - Case 1: random errors
 - Case 2: bursts of 1000 errors
 - Case 3: real-time application (teleconference)



Case 1 – Random Errors

BER: 1 in 1000

1. Assume bit errors are random
 - Messages have 0 or maybe 1 error
- Error correction:
 - Need ~10 check bits per message
 - Overhead: 10 bits
- Error detection:
 - Need ~1 check bits per message plus 1000 bit retransmission ~63.2% of the time
 - Overhead: $1 + 1000 * 0.632 = 633$ bits



Case 2 – Bursty Errors

BER: 1 in 1000

2. Assume errors come in bursts of 1000
 - Only 1 or 2 messages in 1000 have errors
- Error correction:
 - Need $\gg 1000$ check bits per message
 - Overhead: > 1000 bits
- Error detection:
 - Need 32? check bits per message plus 1000 bit resend 2/1000 of the time
 - Overhead: $32 + 1000/1000 * 2 = 34$ bits



Case 3 – Teleconference App

3. Assume bit errors are random

BER: 1 in 1000

- Error correction:
 - Need ~10 check bits per message; Overhead: 10 bits
- Error detection:
 - Need ~1 check bits per message plus 1000 bit retransmission 63.2% of the time; Overhead: $1 + 1000 \cdot 0.632 = 633$ bits
- Solution: Neither!
 - No error detection/correction.
 - Reduction of latency is critical for the app.
 - If a 32 x 32 pixel image (~1000 bits)... can you afford having an error in a single pixel?
 - YES!



Error Correction in Practice

- Heavily used in physical layer
 - LDPC is the future, used for demanding links like 802.11, DVB, WiMAX, LTE, power-line, ...
 - Convolutional codes widely used in practice
- Error detection (w/retransmission) is used in the link layer and above residual errors
- Correction also used in the application layer
 - Called Forward Error Correction (FEC)
 - Normally with an erasure error model
 - E.g., Reed-Solomon (CDs, DVDs, etc.)



Key Concepts

- Framing allows us to send full messages instead of bits
- Redundant bits are added to messages to protect against transmission errors
- Two recovery strategies are retransmissions (ARQ) and error correcting codes (FEC)
- The Hamming distance tells us how much error can safely be tolerated
- The optimal recovery strategy depends on the error expected in the channel and the app requirements

