

CSE160: Computer Networks

Project 1 -- High Level Design

2020-09-08



**Professor
Alberto E. Cerpa**



Design Choices

- In the following slides I propose a **high level design** scheme for project 1
- **HOWEVER**, this is not the **ONLY** way you can design project 1
- The bottom line is that whatever it is your design, you should try to work on:
 - **Scalability**: modular design with clean interfaces → easy to use by other modules
 - **Efficiency**: reduce the overhead as much as possible (not at the expense of scalability)
- Ultimately, **you** are in **control** of your code



Flooding Design

- The cleanest design involves having a module and component just for **flooding**
- This can be used by other applications whenever they need to use the service (**Hint: project 2!!**)
- This means having your own flooding header information that should include:
 - **Source addr** of the node initiating the flood
 - Monotonically increasing **sequence number** to uniquely identify the packet
 - A **TTL** field to avoid the packet looping forever



Flooding Design (2)

- It is important to know from **which neighbor** you receive a packet
- I would recommend adding a **Link Layer** module with **source** and **destination addresses**
 - This changes hop by hop as the packet get flooded over the network (see slide 9)
- Having a **node table** (one entry per node) should be useful to implement a cache
 - The cache contains the largest sequence number seen from any node's flood




Flooding Design (3)

- To access the payload of any packet your header structures should end with a **zero length array**
 - Please use the trick/method described in **slide 27** (Why a zero length array?) of the **C Tutorial** (Files → Projects → C-Tutorial.pdf)
- Two options for local level addressing
 - **Wireless**: broadcast medium, so you should take advantage and broadcast to send 1 packet only
 - **Wired**: you should forward packets to each of your neighbors serially




Ideal Flooding



- Only **one flood** packet **per link**
- Nodes send packets on **each interface** to **each neighbor** (do not broadcast to all interfaces)
- Nodes use the **cache** to check for **duplicates**
 - F does NOT send two flood packets to G
 - F does NOT send C's flood packets back to D
 - F does NOT send D's flood packet back to C
 - G does NOT send any flood packet to neither E or F




Wireless Flooding using Broadcast



- Two flood packet per link (double overhead of ideal case)
- Sending nodes must check the cache
 - All flood packets they sent are being received back
 - Each node must constantly check the cache to avoid sending duplicates



Wired Flooding



- Only **one flood** packet **per link** (like ideal case)
- Nodes send packets on **each interface** to **each individual neighbor** (do not broadcast)
- Each node **MUST** know their neighbors (Neigh. Discovery)
 - Send one flood packet to each neighbor
 - Do NOT send to the neigh you received the flood packet from
 - Use the link layer interface to figure out which neighbor sent the packet
- Nodes use the **cache** to check for **duplicates**




Putting it all together

Application Payload

Flooding Header

Link Layer Header

Payload		
Flood Source	Sequence Number	Time to Live
Source Address		Destination Address



APP

FH

LLH

Hello		
A	120	15
A		B

If NOT in the cache (check FH source), forward

Hello		
A	120	14
B		C

- B checks the cache, and forwards the flood message to other neighbors (not A) → implies we know neighbors (Neigh. Discovery)
- Note that the Link Layer header src and dst addresses change hop-by-hop
- Also note that the Flooding header does NOT change, except for the TTL, which is decreased at each hop



Putting it all together (2)

- Note that C node can figure out the **flood's source** since **it is in the FH**
 - This is the index to the Flooding cache (source node address)
- If the Flood sequence number is **NOT** in the cache (or **larger** than one in the cache), we do the following:
 - Update the cache with a new value
 - Forward the message
- The **Link Layer** knows the neighbor that sent the message
 - LL must provide this info when passing the payload to modules
 - We avoid sending the flood message back to node that sent it
- It also knows which **other neighbors** to send the messages (we assume that Neighbor Discovery is running)
- The logic above **slightly changes** if you are **broadcasting messages** at each hop



Neighbor Discovery Design

- Design involves having a module and component just for **neighbor discovery**
- This can be used by other applications whenever they need to use the service (**Hint: flooding and project 2!!**)
- This means having your own flooding header information that should include:
 - **Request** or **Reply** field
 - Monotonically increasing **sequence number** to uniquely identify the packet
 - **Source address** information can be obtained from the link layer



Neighbor Discovery Design (2)

- You need to send neighbor discovery packets **periodically**
 - Nodes could **die** at any time
 - The **quality** of the links might **degrade**
 - Use a **timer** and post a task to do this periodically
 - What is a good timer to avoid overloading the network?
- Upon reception of a neighbor discovery packet, the receiving node must **reply** back
- The mechanics are very similar to Ping and Ping Reply, you could copy or reuse the code
- When getting a reply back, the node should gather **statistics**



When is a node a neighbor?

- If a neighbor node replies to **1 out of 100** neigh discovery packets, is that node a neighbor?
- If a neighbor node has **failed** to reply to the **last 5** neigh discovery packets, is that node active?
- Each neighbor discovery module should gather **statistics** about what is the **quality of the link** to each of its neighbors
 - E.g. responded to 10 out of 20 packets → 50% link
- It should also have **thresholds** to determine when a neighbor node is **no longer responding**
 - E.g. failed to respond to the last 5 consecutive packets, neighbor might be dead or link is gone



How to gather statistics?

- TinyOS has a **statically** allocated **data model**
 - Not possible to dynamically allocate memory
- In this context, the **simplest** thing to implement is a **running average** to gather link quality stats
 - We have two variables per neighbor:
 - o X: total packet received
 - o Y: total packet sent
 - Each neighbor discovery packet header has a monotonically increasing sequence number
 - If a skip in the sequence, it means loss packet(s)
 - o E.g.: last seq. received: 116, next seq.: 120 → this means three losses, 117, 118, and 119
 - o $X(t) = 59 \rightarrow X(t+1) = 60$; $Y(t) = 116 \rightarrow Y(t+1) = 120$
 - o Link quality $(t+1) = X(t+1) / Y(t+1) = 60 / 120 = 50\%$



Neighbor Table

- You should have a neighbor table with three (3) fields:
 - Neighbor address
 - Quality of the link (percentage)
 - Active neighbor (yes/no)
- You may also **remove** from table neighbor that are **no longer active**
 - Based on the previously defined threshold
- Also note that the neighbor table has **VERY** important information for a lot of applications
 - Make your design such that other components can **query** the content of the neighbor table
 - **Routing (project 2!!)** will need to **access** the content of this table




Putting it all together

Neigh Discovery Header

Link Layer Header

reQuest or Reply	Sequence Number
Source Address	Destination Address



- When **timer** fires off, B **broadcasts** (Br) a neighbor **Request** (Q) discovery packet **increasing** the **last sent sequence number**
- Each **neighbor** node that **receives** a **Request** (Q), sends back a **Reply** (R) immediately
- Upon reception of **R** messages, B **updates** the **neighbor table** with the new **neighbor info** and **statistics**

