

CSE160: Computer Networks

Project 2 – LS Routing High Level Design

2020-10-02



**Professor
Alberto E. Cerpa**



Design Choices

- In the following slides, I propose a **high-level design** scheme for project 2
- **HOWEVER**, this is **NOT THE ONLY** way you can design project 2
- The bottom line is that whatever it is your design, you should try to work on:
 - **Scalability**: modular design with clean interfaces → easy to use by other modules
 - **Efficiency**: reduce the overhead as much as possible (not at the expense of scalability)
- Ultimately, **you** are in **control** of your code



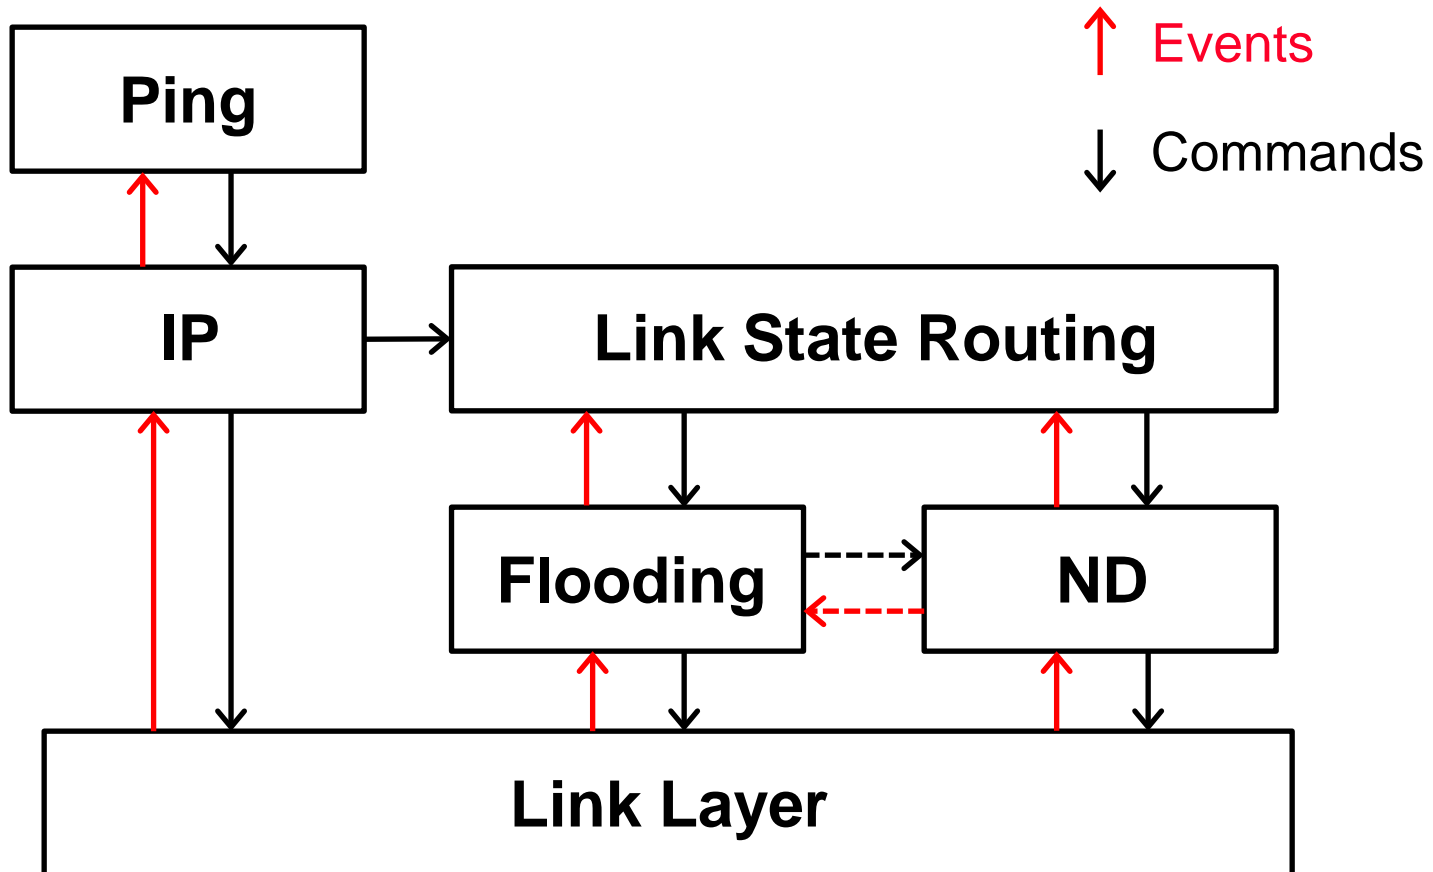
Routing Design

- The cleanest design involves having a module just for **Link-State Routing (LS)**
- In addition, you will need to develop an **IP** module (like the one suggested for the Link Layer in project one)
- The **IP** module will do all the **data forwarding**, using the routing table information provided by **LS**
- This **IP** module can be used by other applications whenever they need to use the **IP service** (**Hint: project 3!!**)
- In this project, we can **test** the IP data forwarding functionality by using **Ping** on top of the IP service
- The **LS** module will oversee the **update** of the **Routing Table** used by the **IP** module



Routing Design (2)

- In addition, you will need both a **Neighbor Discovery (ND)** and **Flooding (F)** modules with appropriate interfaces to hook everything up
- The whole stack could look like this:



LS Advertisement (LSA) Packet Format

- The LS advertisement (LSA) packet should include:
 - **Source Address** of the node sending the info
 - This information may be obtained from the Flooding Header if passed to the LS Event Handler (omitted in LSA then)
 - Monotonically increasing **sequence number** to uniquely identify the LSA update
 - A **list** of tuples with the following format:
 - Tuple_n : $\langle \text{Neighbor Address}_n, \text{Cost}_n \rangle$
 - List: $[\langle \text{Neigh Addr}_1, \text{Cost}_1 \rangle, [\langle \text{Neigh Addr}_2, \text{Cost}_2 \rangle, \dots]]$
 - You could also add other fields like number of tuples in the list, or any other thing you may deem necessary
- To access the payload at any layer, your header structures should end with a **zero-length array**
 - Please use the trick/method described in **slide 27** (Why a zero-length array?) of the **C Tutorial** (Files → Projects → C-Tutorial.pdf)



LSA Update Cache

- It may be useful to maintain an **LSA update cache** in your design, keeping track of the **last sequence number** for any advertisement sent by any node
- This is useful to discard old LSA updates with old (stale) information
 - If your node receives an **LSA update** from a node, with an **older sequence number than the one in your LSA update cache**, it gets **discarded**
 - This means that **no** new shortest path computation will be run [see later]
- The structure could be like the flooding cache you already use in the Flooding module
 - You need **one entry per node** in the topology in the **cache** (like Flooding)



Bootstrapping

- Initially, you **don't** want to **start** the Dijkstra's SP calculation **until** you receive **info** from **all the nodes** in the system
 - This allows the node to calculate SP to all destinations **once you have complete topological knowledge**
- The timing sequence should be roughly:
 - ND starts immediately, learning the list of local neighbors
 - If Flooding **depends** on ND, the node **needs** to wait until ND discovers all neighbors
 - This can be done, for example, by posting a task that periodically checks if the neighbors have changed
 - The Flooding event handler can keep the state whenever ND signals for a neighbor update (e.g. increase the size of a variable)
 - If after the timeout the state has changed, it reposts itself; if not, it starts the process
 - Since LS **depends** on ND, we can use a similar process as above
 - LS **must also wait** to receive **all** the initial LSAs from **all the nodes**
 - Use a timer before starting the process



Initial SP Computation

- Once the node has received the initial LSA updates from **all nodes** in the system, the node has **all the info** to start computing SP to all nodes
- The computation is run **once** at the beginning, and then **after** any **update** [see later]
- For ideas on how to implement **Dijkstra**, look at **Peterson Chapter 3.3** and the class **Lecture 09**
 - You could also use your **CSE 100 code** that asked you to implement **Dijkstra** (I think I assigned it when I taught it)
 - Remember to save not only the **Cost**, but also the **Next Hop** information (it would be a good idea to also save a few alternative paths too, e.g. **top shortest**, **2nd** and **3rd shortest paths**)
- The output of the algorithm, should be used to update the Routing Table...



Routing Table

- This is the **main data structure** that is **used** for **data forwarding**
- It looks something like this:

Destination	Next Hop	Cost	Backup Next Hop	Backup Cost
3	5	7	4	9
9	11	5	10	14
...

- You can decide what data structure to use (your choice)
- Your LS module should have **commands** to **access** the content of the **table** by other modules
 - This is used by the **IP** module to perform **data forwarding**



Adaptation to Dynamics

- When the **ND module** discovers a **change** in the **neighbors list**, it should **signal** to **all** the other **modules** wired to it
 - Changes can be:
 - **Adding** a new **node**
 - **Removing** an old unresponsive **node**
 - **Changing** the **cost** of an existing **node**
- Note that the **ND module** should be implemented to **signal** an event to the **LS's event handler** every time there is a **change** in the **Neighbor List**
 - This should have been already implemented if your **Flooding** module was **wired** to the **Neighbor Discovery** module for proj 1
- This signal triggers **two actions** (called **triggered update**):
 1. A **new LSA** should be **flooded** to all nodes in the system
 2. A **new SP** computation should be **done** [see next]



SP Computation in Steady State

- After the node has run the initial SP computation, the node should only run the SP computation every time there is a change in the topology
- So, in steady state, the SP computation is run every time that:
 - The node receives a new LSA update from another node
 - The local ND module changes the neighbors list (by id or cost of the node) [see previous slide]
 - The points above are the ones that trigger a change in the topology!



IP Module

- The **IP** module oversees doing **data forwarding**
- The **IP** module consults the **routing table** with the list of **shortest paths** to **every destination** maintained by the **LS** routing module
- The IP header has the following fields:
 - **Source** address, **Destination** address, and a **TTL** field
 - The header could have an optional **Protocol** field (like the IPv4's Protocol field and the LL Type field, see next slide)
 - This is **NOT** absolutely required, since you are going to use **only** Ping in this project, and next project you will wire **only** a TCP module
 - If you ever want to wire **multiple modules concurrently**, then this field is required for (de)multiplexing
 - This could be very useful for live debugging, so I would add it
 - E.g.: sending pings directly over IP (no reliability) and then over TCP (with reliability) and check the difference



Data Forwarding

- The **forwarding rules** are as follows:
 - If a packet has an **IP destination address** that is **NOT** the **current node**, then the packet must be **forwarded**
 - The **next hop** is determined by consulting the **routing table** maintained by the **LS routing** module
 - If the packet has an **IP destination address** that **matches** the **current node**, then the packet is **delivered** to the **local module** expecting the packet
 - In this latter case, the packet is **NOT** forwarded to **any other node**
- Flooding and ND Packets:
 - In this design, Flooding and ND packets are sent directly over the Link layer, so they should **NOT** be carried by the **IP** layer
 - In the explanation of project 2 (in the description), flooding is somehow integrated with data forwarding, but I find this may be confusing for many students




Example: LSA Distribution

LSA Packet

Flooding Header

Link Layer Header

Source	Seq Num	Num Entries	Node	Cost
Flood Source	Sequence Number		Time to Live	
Source Address		Destination Address	LL Type	



LSA

A	27	1	B	1
---	----	---	---	---

FH

A	52	15
---	----	----

LLH

A	B	Flood
---	---	-------

A	27	1	B	1
---	----	---	---	---

A	52	14
---	----	----

B	C	Flood
---	---	-------

- Flooding works as before like in project 1 (no explanation)
- Notice that we added a **Link Layer Type** in the LL Header
 - This is because the Link Layer is **multiplexing** connections from the **IP**, **Flooding** and **Neighbor Discovery** modules in our design → need **identifier** in the **header** to know which **module to signal an event** when packet **arrives** at the LL module
 - This is like the **Ethertype** field in the **Ethernet Header**



Example: LSA Distribution (2)

- The Link Layer **source** and **destination** fields change **hop by hop**
- The **only** field that **changes hop by hop** in the **Flooding Header** is the **TTL**
- The **LSA Packet** **never changes** hop by hop
 - It should be forwarded by Flooding **AS IS**
- Please note that in this example, node **A** has **only one** neighbor (B), so the list in the LSA is very short (just **one tuple**)
 - In your design, you should **NOT** send just **one tuple at a time**, you want to send **as many as you can fit in a single packet**
 - In TinyOS, by default, the maximum packet size is **29 bytes**, but there is some overhead at the lowest levels of the stack
 - This means that for nodes with **lots of neighbors**, you may need to send (flood) **multiple LSA** packets
 - Consult the **TOS Packet Format** in the code




Example: Data Forwarding

Ping Packet

IP Header

Link Layer Header

Ping Payload		
Source Address	Destination Address	Time to Live
Source Address	Destination Address	LL Type



Ping

IP

LL

Ping Payload		
A	C	15
A	B	IP

Ping Payload		
A	C	14
B	C	IP

- A sends a ping to C
- C is **NOT** directly connected to A, i.e. multi-hop routing is required
- B forwards the IP packet to its destination



Example: Data Forwarding (2)

- Data forwarding as the packet traverses the network:
 - Node A's IP module receives a Ping payload from the wired Ping module, with destination IP address C
 - Node A's IP module consults the routing table for destination C destination, and finds that next hop is B
 - Node A's IP module assembles a new IP packet with source, destination addresses and TTL (A, C, 15) and the ping payload
 - Node A's IP module sends the packet to the LL module, with LL destination address B, i.e. the next hop from the routing table
 - Node B's IP module receives the IP packet from the LL module
 - Since the IP destination address is NOT B, then B consults the routing table and forwards the packet to C (next hop)
 - Node C's IP module receives the IP packet from the LL module
 - Since the IP destination address matches C, it delivers the packet to Ping



Food for Thought

- The current design assumes **perfect knowledge** and **no errors**
 - What happens if you get errors?
- A more **resilient** design might be to have LSAs **expire** after a certain amount of time
- This design involves **periodically** sending LSAs **before** they **expire**, even if **nothing** changes in the topology, to **renew** the LSAs
 - E.g., OSPF's LSAs **expire** after **1 hour**, and they are **renewed** every **30 minutes** (**50 minutes** with flooding reduction)
 - A design like this involves getting a **timer** to send the LSAs **periodically**, in addition to **triggered updates**

